

Artificial Stupidity Research Dossier

VERSION 0.9.4 • CLAIM PRECISION EDITION

A MONAHINGA™ EVIDENCE PROJECT

Raymond Anthony Gomez | September 15, 2026



Intelligence proposes. A separate gate challenges. Evidence informs the decision. A human retains final authority.

CURRENT STATUS Working research prototype. Human external review submitted; findings pending. Independent reproduction has not been completed.

Status and purpose

Artificial Stupidity is a working research prototype that adds a separate acceptance boundary around an AI research loop. It asks a narrow but consequential question: when an AI produces a better metric, what evidence and authority should be required before anyone accepts the change?

The gate applies declared rules and returns **KEEP**, **REJECT**, or **ESCALATE**. It preserves the machine decision separately from human authorization. The public record includes 74 automated CPU tests, frozen adversarial cases, a six-run H100 experiment, raw logs, source snapshots, SHA-256 manifests, a disclosed execution failure and repair, a same-team portability exercise, a live demonstration, and a frozen external-reviewer packet.

The project demonstrates an executable evidence-gating pattern. It does not establish general model improvement, causal benefit from smaller batch size, independent reproduction, peer review, production readiness, or an AI-safety result. A corrected packet has been submitted for independent human review; findings remain pending.

Executive finding

Artificial Stupidity began as a joke after an AI comparison favored Andrej Karpathy's work. The joke led to a historical search, then to a working extension of Karpathy's `autoresearch` loop. The extension does not replace the optimizer. It adds a decision boundary outside the optimizer's editable training file.

The most important project result is organizational rather than causal: a model-training candidate recorded better metrics, yet the machine still stopped because it lacked authority to certify human understanding. A named person reviewed and authorized the bounded decision separately.

The project's own interpretation then failed a different test. Claude, acting as an AI adversarial pre-reviewer, identified a fixed-wall-clock confound and several phrases that could be read more broadly than the evidence allowed. Version 0.9.4 narrowed the public claims without altering the measurements. That correction strengthens the central principle: an apparently persuasive result must survive challenge, and the challenge record should remain visible.

Evidence labels used in this dossier

Label	Meaning
Recorded	Directly supported by a preserved project artifact, public repository state, or cited source.
Interpretation	A reasoned synthesis that should not be mistaken for a measurement.
Proposal	A design or next step that has not yet been executed.
Unknown	A material fact the current record cannot establish.

Selected intellectual roots

The phrase “artificial stupidity” and the concern behind it predate this project. The lineage is fragmented rather than one continuous discipline. This section identifies selected roots; it does not claim an exhaustive historical search.

Period	Selected root	Relevance and boundary
1950	Alan Turing discussed deliberate mistakes in the imitation game.	Connects machine fallibility with appearing human; not an evidence-governance proposal.
1991–1992	Early conversational competitions and reporting used intentional limitation to make programs appear believable.	Shows artificial imperfection as a design tactic; the 1992 article is historical context, not a scientific foundation for this project.
2003–2004	Lars Lidén described intentional mistakes in game AI; Peter and Dorothy Denning used institutional satire to distinguish calculation from judgment.	Shows that optimal machine performance and good human outcomes are not identical.
2016	*Concrete Problems in AI Safety* described failures such as reward hacking, harmful side effects, and distributional change.	Supports the general concern that a score can improve while the intended purpose is missed.
2018	Trazzi and Yampolskiy proposed safer AGI through intentional limitations on capability.	A genuine “artificial stupidity” safety precursor, but it changes the intelligent system; this project instead changes the acceptance boundary.
2019	Selective prediction formalized abstention as a useful outcome.	Provides an ancestor of ESCALATE : uncertainty can become visible action rather than fabricated certainty.
2021–2024	Work on judgment, self-correction, and sycophancy challenged the assumption that the same model can always provide a trustworthy second opinion.	Supports separation and external evidence; does not prove that every form of self-review fails.
2025	Andrej Karpathy coined “vibe coding.”	Names the builder practice that made rapid, conversational construction central to this project.
2026	Karpathy released <i>autoresearch</i> , a compact autonomous experiment loop.	Supplied the MIT-licensed engine sampled and extended by this project.

The defensible novelty claim is narrow. The project did not invent criticism, holdouts, abstention, human oversight, the phrase “artificial stupidity,” or autonomous experimentation. It combines these ideas in a public prototype that separates proposal, evidence, gate decision, challenge, and human authorization.

Origin story

The comparison

In a separate conversation, Claude told Ray that Andrej Karpathy was a better vibe coder and pointed to Karpathy's `autoresearch` repository as a strong example. Ray took the comparison personally—but productively. He downloaded the model and asked what consequential decision the upstream loop left unresolved.

`autoresearch` changes model-training code, runs a time-bounded experiment, measures validation bits per byte, and keeps or discards the candidate. It makes rapid autonomous discovery concrete. It also creates a useful design question: should the system proposing and measuring a change hold sole authority to accept it?

The joke

A Harry Potter exchange supplied the phrase “Artificial Stupidity.” The historical search showed that the phrase had bones. Earlier meanings often involved deliberate errors, humanlike imperfection, capability restriction, institutional foolishness, or the difference between intelligence and judgment.

This project chose a non-degrading interpretation:

Do not make the intelligent engine worse. Make acceptance harder to earn.

The build

The upstream optimizer remained credited and materially intact. A separate deterministic gate was placed outside its editable training file. The first defined build, testing, hardening, evidence packaging, public demonstration, and initial external-review submission occurred over approximately 55 hours from the precursor conversation. That is a creator-reported build chronology, not evidence of scientific validity.

System architecture

Roles

Role	Responsibility	Authority limit
Proposer	Suggest a candidate change and state the intended result.	Cannot approve its work merely because a score improved.
Evidence harness	Run declared comparisons and preserve measurements, logs, source, environment, and hashes.	Cannot turn a correlation into a causal explanation by itself.
Gate	Apply declared rules and return KEEP, REJECT, or ESCALATE.	Cannot exceed delegated authority or certify human understanding.
AI challenger	Attack assumptions, wording, logic, missing controls, and failure paths.	Does not count as independent reproduction or human review.

Role	Responsibility	Authority limit
Human authorizer	Accept responsibility for a bounded decision after inspecting the record.	Cannot convert limited evidence into a universal finding.
External reviewer	Reproduce declared procedures and publish findings under a disclosed relationship.	Must not be described as having validated the project before completing the work.

Gate outcomes

Action	Decision meaning	Example condition
KEEP	The candidate cleared the machine rules delegated to the gate.	Required evidence is present and no declared disqualifier fires.
REJECT	The candidate hit a declared failure condition.	The gain is too weak, a critical failure exists, or a resource boundary is exceeded.
ESCALATE	The result may be promising, but evidence or machine authority is insufficient.	Required evidence is unknown, conflicting, or reserved for human judgment.

Evaluated evidence

The current model-training gate can inspect whether the primary metric improved, whether the result repeated, whether it survived a separately stored holdout, whether resource use remained within a configured boundary, whether critical failures were recorded, whether the claim is falsifiable, whether the change remained within declared complexity limits, and whether a responsible human has reviewed what the machine cannot certify.

Malformed, impossible, or missing required evidence fails closed under the declared cases. This behavior is demonstrated by deterministic tests. It is not a guarantee against every attacker, implementation flaw, or real-world failure.

Development and evidence chronology

Stage	What happened	What the stage established
Origin	A bruised ego and a Harry Potter joke prompted the historical and technical question.	A human reason to investigate; no empirical claim.
Prototype	A separate KEEP / REJECT / ESCALATE gate was added around autoresearch.	The architecture could be made executable.

Stage	What happened	What the stage established
First H100 run	A one-line candidate improved ordinary validation, but no separately stored holdout existed.	The gate could withhold automatic acceptance after an encouraging metric.
Six-run H100 record	Three baseline and three candidate runs were interleaved under a fixed-time budget.	A small bounded measured difference in this setup; not causal attribution.
Human authorization	The machine escalated on intentionally unknown human comprehensibility; Ray separately authorized KEEP.	The machine and human decision records could remain distinct.
Late import failure	Final gate import failed after runs and evaluations completed; preserved inputs allowed reconstruction and repair.	Evidence packaging supported diagnosis after automation failed.
Adversarial v0.9.1 campaign	Eight of nine attack probes produced unjustified KEEP outcomes.	The first boundary was materially weak.
v0.9.2 repair	Tested attack paths were repaired and the frozen follow-up matrix passed.	The identified weaknesses were addressed under the declared cases.
Bill X-Ray exercise	Three same-team, pre-labeled legislative claims produced the expected decisions.	Initial portability of the decision pattern; not independent validation.
Claude challenge	An AI critique attacked the claims and found the fixed-time confound and interpretive overreach.	The public explanation needed correction.
v0.9.4 precision pass	Claims were narrowed while measurements and prior releases remained preserved.	A synchronized, more defensible claim boundary.
Human external review	The corrected packet was submitted to an outside reviewer.	A review process began; no findings yet.

H100 experimental record

Candidate and protocol

The candidate changed one functional line:

TOTAL_BATCH_SIZE = 2**19 baseline

TOTAL_BATCH_SIZE = 2**18 candidate

Three baseline and three candidate runs were interleaved:

B1 / C1 / B2 / C2 / B3 / C3

The runs used the same disclosed H100 SXM environment and a fixed wall-clock training budget. Ordinary validation and a separately stored shard were evaluated. The public evidence contains the raw logs, environment record, source snapshot, hashes, structured results, machine decision, repair provenance, and human authorization.

Recorded values

Metric	Baseline runs	Candidate runs	Baseline median	Candidate median	Recorded difference
Validation BPB	0.999791, 1.000213, 1.001808	0.990674, 0.990934, 0.991104	1.000213	0.990934	0.928% lower
Sequestered-holdout BPB	0.998995, 0.999357, 1.001117	0.989901, 0.990126, 0.990284	0.999357	0.990126	0.924% lower
Peak VRAM	—	—	45,060.2 MB	44,908.2 MB	0.337% lower

Lower BPB is better. Each of the three candidate runs recorded lower validation BPB than each of the three baseline runs. The same ordering appeared on the separately stored holdout. Those three-by-three relationships create nine all-by-all cross-arm comparisons on each measure. They are not nine independent trials.

What the record supports

Under this one disclosed fixed-time setup, the preserved candidate runs recorded lower validation and sequestered-holdout BPB than the preserved baseline runs. The candidate also used slightly less peak VRAM in the summarized record.

The fixed-time confound

The protocol held wall-clock time constant but did not hold optimizer steps or tokens processed constant. Changing batch size may have changed how many updates or tokens each arm completed. Batch size was therefore bundled with untracked training exposure.

This prevents a causal conclusion that halving batch size produced the improvement. It also prevents general claims about other models, datasets, hardware, or training budgets. The recorded comparison remains useful as a bounded engineering observation under its own procedure.

Decision chain

The empirical fields satisfied the declared machine-level boundary, but `human_comprehensibility` was intentionally unknown. The machine was not authorized to claim that a person understood the change. It returned **ESCALATE**.

Ray reviewed the one-line change and bounded record and separately authorized **KEEP** on September 13, 2026. This was author authorization, not an independent technical finding. Both records remain preserved.

The central demonstration is the stop itself: the system reached the edge of its delegated authority, exposed the unresolved condition, and left a separate record for a named human decision.

Failures and corrections

Late execution failure

After the six training runs and evaluations completed, the runner encountered `ModuleNotFoundError: No module named 'gate'` while reaching the final decision step. The original run therefore did not exit successfully as an end-to-end process.

The archive had already preserved the raw logs, environment, source snapshot, hashes, and structured partial results. The decision was reconstructed locally from those captured inputs without additional training. The import path was repaired and covered by a regression test. The original failure remains public.

Recorded: evidence preservation supported reconstruction and diagnosis.

Not established: production resilience or immunity from future execution failures.

Adversarial trust-boundary failure

The first frozen adversarial campaign found a more serious logical weakness. Eight of nine attack probes produced unjustified **KEEP** outcomes in v0.9.1. These cases included evidence-provenance and trust-boundary problems that ordinary happy-path tests had not exposed.

Version 0.9.2 repaired the tested paths, added explicit handling for unknown states and remedies, and passed the frozen follow-up attack and control matrix. Both receipts remain public.

Recorded: the declared attacks found weaknesses, and the follow-up version passed those same frozen cases.

Not established: comprehensive security, cryptographic identity, or protection against a party controlling policy, evidence, and runtime.

Bill X-Ray portability exercise

Three Illinois legislative-claim cases were selected, labeled, and frozen by the same project team before a separate adapter was implemented:

Case	Predeclared action	Reason
Supported statutory statement	KEEP	The cited statutory language supported the bounded claim.
Mandate-and-funding overclaim	REJECT	The source contradicted the claim that adoption was mandatory or funding guaranteed.
Contractor-outcome forecast	ESCALATE	The statute did not establish future win rate or market demand.

The adapter produced all three expected actions, and the associated portability and security tests passed. This shows that the decision doctrine can be represented outside model training under known cases.

It remains a same-team, deterministic, pre-labeled exercise. It is not legal advice, independent domain validation, an error-rate estimate, or evidence of real-world generalization.

Claude adversarial pre-review

Claude was given the reviewer materials, repository, live demonstration, portfolio site, and profile framing and was asked to inspect the evidence rather than accept the project's narrative. The critique separated what the project genuinely demonstrated from what remained unproven.

The most consequential finding was the fixed-wall-clock confound. Claude also challenged:

- language that made the gate sound institutionally or operationally independent;
- “9/9 wins” without explaining the six underlying runs;
- “second-domain validation” for a same-team pre-labeled exercise;
- a submitted external-review request being read as completed review;
- marketing language that outpaced the evidence.

The critique was not itself independent reproduction, peer review, or human review. It functioned as an internal AI challenger. Accepted corrections were checked against the preserved repository evidence and recorded in the v0.9.4 claim-precision release.

This episode created a useful before-and-after receipt: the project applied its philosophy to its own public claims.

Current claim ledger

Claim	Status	Safe interpretation
The gate runs and returns KEEP, REJECT, or ESCALATE.	Supported	Demonstrated in public code and deterministic tests.
The gate can stop automatic acceptance after a metric improves.	Supported	Demonstrated by the preserved decision chain.
Machine judgment and human authorization are stored separately.	Supported	Demonstrated in the public record.
The candidate recorded lower BPB under the six-run fixed-time setup.	Supported and bounded	Applies only to the preserved runs and disclosed procedure.
Halving batch size caused the improvement.	Unsupported	Steps and tokens were not held constant.
The result generalizes across models, data, hardware, or domains.	Unsupported	No such study has been completed.
The Bill X-Ray exercise independently validates portability.	Unsupported	Same-team, deterministic, and pre-labeled.

Claim	Status	Safe interpretation
The gate is an AI-safety system.	Unsupported as a result claim	It is a research prototype demonstrating a gating pattern.
The project has completed independent review or reproduction.	Unsupported	Human review was requested; findings are pending.
The 74-test suite proves real-world reliability.	Unsupported	It shows conformance to declared deterministic cases.

Falsifiable next experiment

The highest-value empirical follow-up is designed to separate the observed engineering result from a causal batch-size claim.

Before any new result is inspected, freeze:

1. Six matched baseline and candidate pairs.
2. Interleaved execution.
3. Matched seeds within each pair where the implementation permits, with different seeds across pairs.
4. Per-run wall time, optimizer steps, tokens processed, BPB, holdout BPB, peak VRAM, exits, and failures.
5. A token-matched control that separates batch size from total training exposure.
6. A fixed stopping rule and disclosure rule for failed or excluded runs.
7. Raw win counts without converting a small sample into a significance claim.
8. Independent recomputation of BPB from logits and recomputation of the gate decision from raw evidence.

The justified conclusion afterward would remain bounded: under a named architecture, dataset, environment, and protocol, the candidate did or did not retain its measured advantage when training exposure was controlled.

This experiment is a proposal. It has not been run.

Research program beyond one experiment

A broader evaluation should compare matched conditions such as AI alone, the same AI with self-critique, a separately prompted critic, a critic with external tools and sequestered evidence, and the full system with structured human authorization.

Useful measures include consequential error, unsupported claims, calibration, appropriate abstention, false vetoes, reviewer accuracy, time, cost, and the rate at which humans defer to a confident but wrong recommendation.

The doctrine should be revised or rejected if separated challenge does not outperform simpler review at matched cost, if criticism mainly blocks correct high-value action, if the proposer can routinely game the gate, or if the gate increases automation bias rather than reducing it.

Public and practical path

Public core

The repository, live demonstration, release packets, claims, failures, repairs, and evidence should remain inspectable. Public challenge is part of the credibility strategy.

Potential organizational pilots

The practical opportunity is not to sell a universal safety layer. It is to identify one bounded workflow where an AI produces recommendations and where the required evidence, disqualifying conditions, escalation authority, and acceptance decision can be declared in advance.

Potential pilot settings include autonomous research, coding agents, analytical workflows, regulated review, procurement, and audit preparation. A responsible pilot would require domain expertise, privacy and access controls, threat modeling, operational testing, and a clear human owner.

Evidence before promotion

The external human review and controlled follow-up experiment are more valuable than additional broad marketing. Outreach should route recipients to the smallest useful asset:

- live demonstration for fast understanding;
- executive summary for a three-minute evaluation;
- this dossier for the full case;
- tagged v0.9.4 packet for hands-on reproduction;
- repository for code, logs, tests, provenance, and version history.

MONAHINGA relationship

Artificial Stupidity is presented as **A MONAHINGA™ Evidence Project**. MONAHINGA is the source identity for evidence-backed investigations that recover scattered records, separate documentation from inference, test rival explanations, preserve receipts, and make conclusions reviewable.

USPTO serial 99613654 was recorded in the prior project materials as live/pending in International Class 041 under Section 1(b). Registration has not issued according to that record. The mark does not authenticate the research or create ownership over facts, methods, historical events, public records, inherited code, or the phrase “Artificial Stupidity.” Current trademark status should be verified before any formal filing or commercial representation.

Conclusion

The strongest conclusion is deliberately modest. Intelligence, evidence, judgment, and authority are different functions. A system that can generate and measure an improvement should not automatically inherit the right to approve it.

Artificial Stupidity is a public working prototype of that separation. Its first evidence record is small, confounded for causal interpretation, and not independently reproduced. Its value lies in making the boundary inspectable, preserving failure, inviting attack, and narrowing its own claims when criticism

earns the correction. The project began with embarrassment and a joke; it became more credible when the claim became smaller.

Current project links

- **Live demonstration:** <https://artificial-stupidity.onrender.com/>
- **Repository:** <https://github.com/rayoonbaco/artificial-stupidity>
- **Frozen v0.9.4 reviewer packet:** <https://github.com/rayoonbaco/artificial-stupidity/releases/tag/v0.9.4-claim-precision>
- **Portfolio:** <https://the-receipts.onrender.com/>

Selected references

- [1] A. M. Turing, “Computing Machinery and Intelligence,” *Mind* 59(236), 1950.
<https://doi.org/10.1093/mind/LIX.236.433>
- [2] H. G. Loebner, “In Response,” The Computer Museum, 1991.
<https://tcm.computerhistory.org/Timeline/LoebnerPrize1991.pdf>
- [3] L. Lidén, “Artificial Stupidity: The Art of Making Intentional Mistakes,” *AI Game Programming Wisdom 2*, 2003.
https://www.liden.cc/lars/WEB/Resume/Papers/2003_AIWisdom.pdf
- [4] P. J. Denning and D. E. Denning, “Artificial Stupidity,” *Communications of the ACM* 47(3), 2004.
<https://doi.org/10.1145/986213.986246>
- [5] D. Amodi et al., “Concrete Problems in AI Safety,” 2016. <https://arxiv.org/abs/1606.06565>
- [6] M. Trazzi and R. V. Yampolskiy, “Building Safer AGI by introducing Artificial Stupidity,” 2018.
<https://arxiv.org/abs/1808.03644>
- [7] Y. Geifman and R. El-Yaniv, “SelectiveNet,” *ICML 2019*. <https://proceedings.mlr.press/v97/geifman19a.html>
- [8] M. Falk, “Artificial Stupidity,” *Interdisciplinary Science Reviews* 46(1–2), 2021.
<https://doi.org/10.1080/03080188.2020.1840219>
- [9] N. Shinn et al., “Reflexion,” 2023. <https://arxiv.org/abs/2303.11366>
- [10] J. Huang et al., “Large Language Models Cannot Self-Correct Reasoning Yet,” *ICLR 2024*.
<https://arxiv.org/abs/2310.01798>
- [11] M. Sharma et al., “Towards Understanding Sycophancy in Language Models,” *ICLR 2024*.
<https://arxiv.org/abs/2310.13548>
- [12] Andrej Karpathy, autoresearch, 2026, MIT License. <https://github.com/karpathy/autoresearch>

Research and authorship note

Concept, framing, project direction, and human authorization: **Raymond Anthony Gomez**. Research synthesis, source checking, drafting, visualization, and prototype implementation were developed through iterative work with ChatGPT. Claude later served as an AI adversarial pre-reviewer whose accepted criticisms informed the v0.9.4 claim-precision pass. Historical statements should be checked against the cited primary sources before formal scholarly publication. Experiment claims remain tied to the preserved repository evidence. This dossier explains that record; it does not replace it.