

Artificial Stupidity

EXECUTIVE SUMMARY

A MONAHINGA™ EVIDENCE PROJECT

Version 0.9.4 claim precision edition

Raymond Anthony Gomez | September 15, 2026



Executive conclusion

Artificial Stupidity is a working research prototype that adds a separate acceptance boundary around an AI research loop. An AI may propose and test an improvement, but a better metric does not automatically earn approval. The gate applies declared evidence rules and returns **KEEP**, **REJECT**, or **ESCALATE**, while preserving machine judgment separately from human authorization.

The project demonstrates a functioning evidence-gating pattern. It does not establish a general model-training improvement, completed independent reproduction, peer review, production readiness, or an AI-safety result. External human review has been requested and remains pending.

What was built

- A separate gate around Andrej Karpathy's MIT-licensed autoresearch loop without replacing the upstream training engine.
- Rules covering repeatability, a separately stored holdout, resource use, critical failures, falsifiability, complexity, and human review.
- Explicit **KEEP**, **REJECT**, and **ESCALATE** outputs with reasons and remedies.
- Separate machine-decision and named human-authorization records.
- 74 automated CPU tests, frozen adversarial cases, complete raw H100 logs, source snapshots, SHA-256 manifests, and preserved failure-and-repair receipts.
- A live, GPU-free demonstration that calls the public Python gate rather than displaying hard-coded results.

Recorded H100 result

The tested candidate changed one functional line by reducing total batch size from 2^{**19} to 2^{**18} . Three baseline and three candidate runs were interleaved as B1 / C1 / B2 / C2 / B3 / C3 under the same fixed wall-clock budget.

Metric	Baseline median	Candidate median	Recorded difference
Validation BPB	1.000213	0.990934	0.928% lower
Sequestered-holdout BPB	0.999357	0.990126	0.924% lower
Peak VRAM	45,060.2 MB	44,908.2 MB	0.337% lower

Each of the three candidate runs recorded lower validation and holdout BPB than each of the three baseline runs. This produces nine all-by-all cross-arm comparisons on each measure; it is not nine independent trials.

Interpretation boundary

The comparison fixed wall-clock time but did not hold optimizer steps or tokens processed constant. Halving batch size may have changed training exposure. The record therefore does not isolate batch size as the cause of the measured difference.

The supported statement is narrower: **under this one disclosed fixed-time setup, the preserved candidate runs recorded lower validation and sequestered-holdout BPB than the preserved baseline runs.**

Why the machine stopped

The gate left human comprehensibility intentionally unknown because the machine was not authorized to certify that a responsible person understood the change. It returned **ESCALATE**. Ray reviewed the one-line change and bounded evidence and separately authorized **KEEP**. This was author authorization—not outside validation. The visible separation of those two decisions is the central demonstration.

What criticism changed

The project improved through two different challenge paths:

1. A frozen adversarial campaign exposed a serious v0.9.1 trust-boundary weakness: eight of nine attack probes produced unjustified **KEEP** outcomes. Version 0.9.2 repaired the tested paths and preserved both the failure and the successful follow-up campaign.
2. Claude later served as an **AI adversarial pre-reviewer** and challenged the interpretation of the H100 record. The critique identified the fixed-time confound, ambiguous “independent” language, misleading prominence of “9/9,” overstatement of the Bill X-Ray exercise, and the difference between submitting a review packet and completing a review. Version 0.9.4 narrowed the public claims without changing the underlying measurements.

Claude did not reproduce the experiment or provide independent human review. The external human reviewer has received the corrected packet; findings are pending.

Why the pattern may matter

As AI systems create experiments, code, analyses, and recommendations faster, trustworthy acceptance becomes a separate problem from generation. A persuasive output may improve the stated score while missing the purpose, evidence, risk, or authority boundary that matters to people.

Artificial Stupidity makes that boundary executable and inspectable. The reusable design pattern is to keep the proposer, evidence harness, gate, challenger, and human authorizer in distinct roles with preserved records. Potential settings include autonomous research, coding agents, analytical workflows, compliance review, procurement, and other consequential processes. Any real deployment would require domain-specific rules, threat modeling, privacy controls, qualified review, and operational testing.

Current evidence status

Established	Not established
Working public gate with three explicit decisions	General or causal benefit from smaller batch size
74 deterministic CPU tests	Real-world error rate or safety performance
Six preserved H100 runs under one disclosed setup	Generalization across models, datasets, hardware, or domains
Disclosed late import failure and repair	Cryptographically secret or adversary-proof holdout
Preserved adversarial failure and corrected follow-up	Independent reproduction, peer review, or certification
Same-team, pre-labeled Bill X-Ray portability exercise	Independent legal or cross-domain validation

Highest-value next experiment

Before inspecting new results, freeze a larger paired protocol with six matched baseline/candidate pairs, interleaved execution, matched seeds where possible, logged wall time, optimizer steps and tokens, a token-matched control, a fixed stopping rule, and independent recomputation of BPB and the gate decision. This experiment is proposed; it has not been run.

Inspect the record

- **Live demonstration:** <https://artificial-stupidity.onrender.com/>
- **Repository:** <https://github.com/rayoonbaco/artificial-stupidity>
- **Frozen v0.9.4 reviewer packet:** <https://github.com/rayoonbaco/artificial-stupidity/releases/tag/v0.9.4-claim-precision>
- **Portfolio:** <https://the-receipts.onrender.com/>

Current status: Working research prototype. Human external review submitted; findings pending. Independent reproduction not completed.